

Embedding Self-Organizing Maps into Neural Networks

Stephen Barnes
stbarnes@stanford.edu





Stanford
University

What is a Self-Organizing Map?

What is a Self-Organizing Map?

A mapping from \mathbb{R}^n to \mathbb{R}^m


Dimensionality reduction



What is a Self-Organizing Map?

“Self-organizing” – mapping learned from data

Mapping can be nonlinear (red), unlike PCA (blue)




$R^2 \rightarrow R^1$


Learning an SOM mapping $\mathbb{R}^n \rightarrow \mathbb{R}^m$

Learning an SOM mapping $\mathbb{R}^n \rightarrow \mathbb{R}^m$


Start with an m -dimensional grid graph



$m=$
1



$m=$
2





$m=$
3

Learning an SOM mapping $\mathbb{R}^n \rightarrow \mathbb{R}^m$

Embed into input space \mathbb{R}^n randomly


Each node n has location l_n in \mathbb{R}^n



Learning an SOM mapping $\mathbb{R}^n \rightarrow \mathbb{R}^m$

Show each point to the SOM





When SOM is shown a point, nearby graph nodes move closer



Over many steps, SOM graph copies input distribution

Learning an SOM mapping $\mathbb{R}^n \rightarrow \mathbb{R}^m$


$$\mathbb{R}^2 \rightarrow \mathbb{R}^1$$



Updating the graph embedding

When shown a sample, the closest graph node is the “winner” - competitive learning


Winning node and nearby nodes move closer




Updating the graph embedding

When shown a sample, the closest graph node is the “winner” - competitive learning

Winning node and nearby nodes move closer





Die-off is exponential in graph distance



Learning rate decays with samples seen

Learning an SOM mapping $\mathbb{R}^n \rightarrow \mathbb{R}^m$

m-dimensional grid graph defines an m-dimensional “graph space”




SOM outputs graph-space coordinates of winning node

Learning an SOM mapping $\mathbb{R}^n \rightarrow \mathbb{R}^m$


SOM translates input point to coordinate of nearest grid point in graph space

SOM “unwinds” 1D manifold in \mathbb{R}^2 to a line in \mathbb{R}^1

Input points in \mathbb{R}^2



Graph-space points in \mathbb{R}^1




History of SOMs



Designed by Prof. Teuvo Kohonen
in the 1980s

History of SOMs

Building on 1970s models from neuroscience
and morphogenesis models from 1950s





Pictured:
1970s neuroscience models

Inserting SOMs into NNs



Neural networks transform manifolds to make categories separable

Maybe SOMs can do this better?



Pictures by Chris Olah
<http://colah.github.io/posts/2014-03-NN-Manifolds-Topology/>


Inserting SOMs into NNs



Inserting SOMs into NNs

Problem: SOMs are non-differentiable! Backpropagation is impossible


Nodes win all input points inside Voronoi cell – piecewise constant output



Inserting SOMs into NNs

Solution: every node wins, but some win more than others

Output is weighted mean of nodes' graph locations,
weights decaying with distance to input point



Average of graph positions

Weights decaying with distance

$$\text{Output}(p) = \frac{\sum_v g_v \exp(-\beta \|p - l_v\|_2)}{\sum_v \exp(-\beta \|p - l_v\|_2)}$$


Output of SOM

Normalized by sum of weights

Inserting SOMs into NNs

Solution: every node wins, but some win more than others

Output is weighted mean of nodes' graph locations,
weights decaying with distance to input point



Graph position of node v Decay parameter

$$\text{Output}(p) = \frac{\sum_v g_v \exp(-\beta \|p - l_v\|_2)}{\sum_v \exp(-\beta \|p - l_v\|_2)}$$

Output of SOM Input point Embedding location of node v

Inserting SOMs into NNs

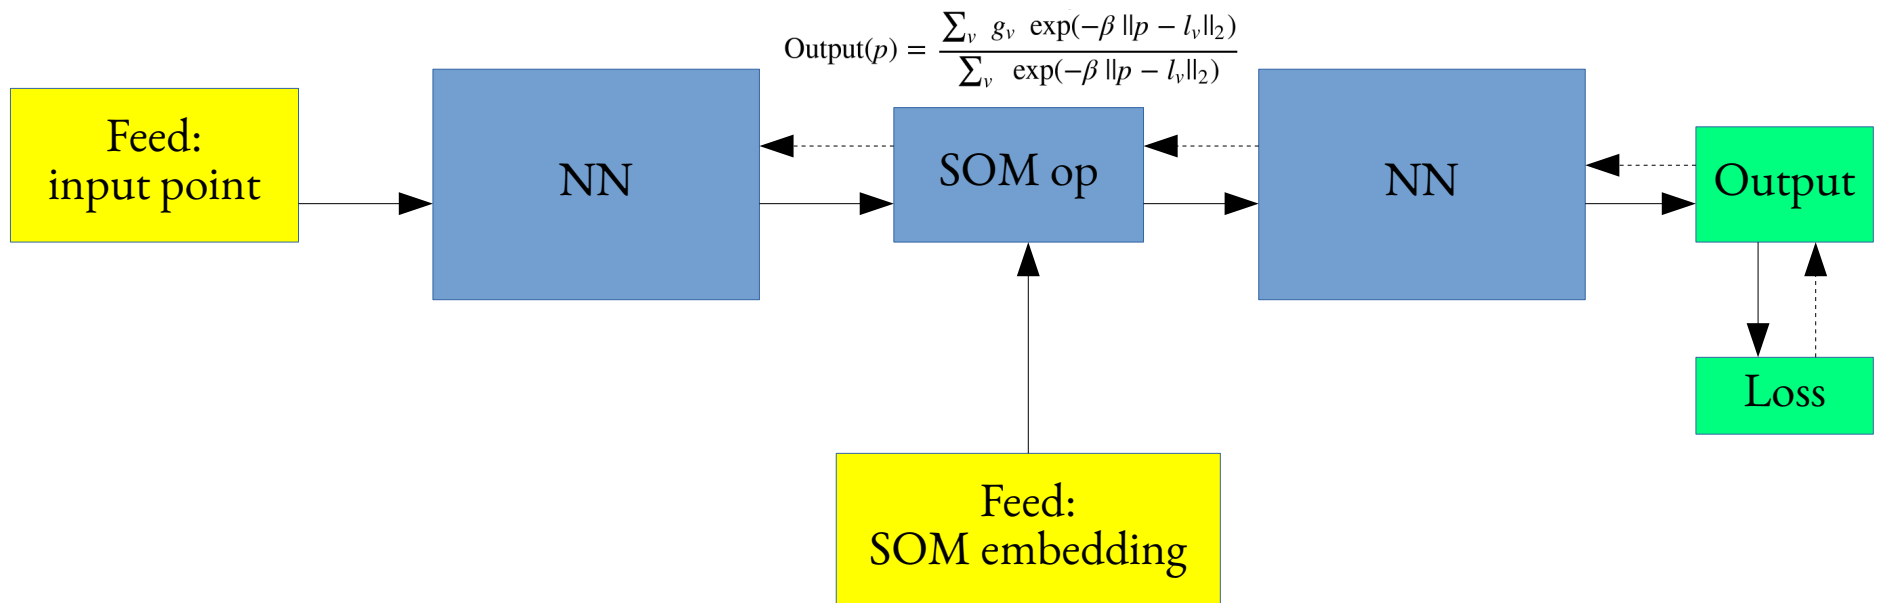
If node embedding is constant,
this is a differentiable function of p .

Derivatives can flow back through SOM!

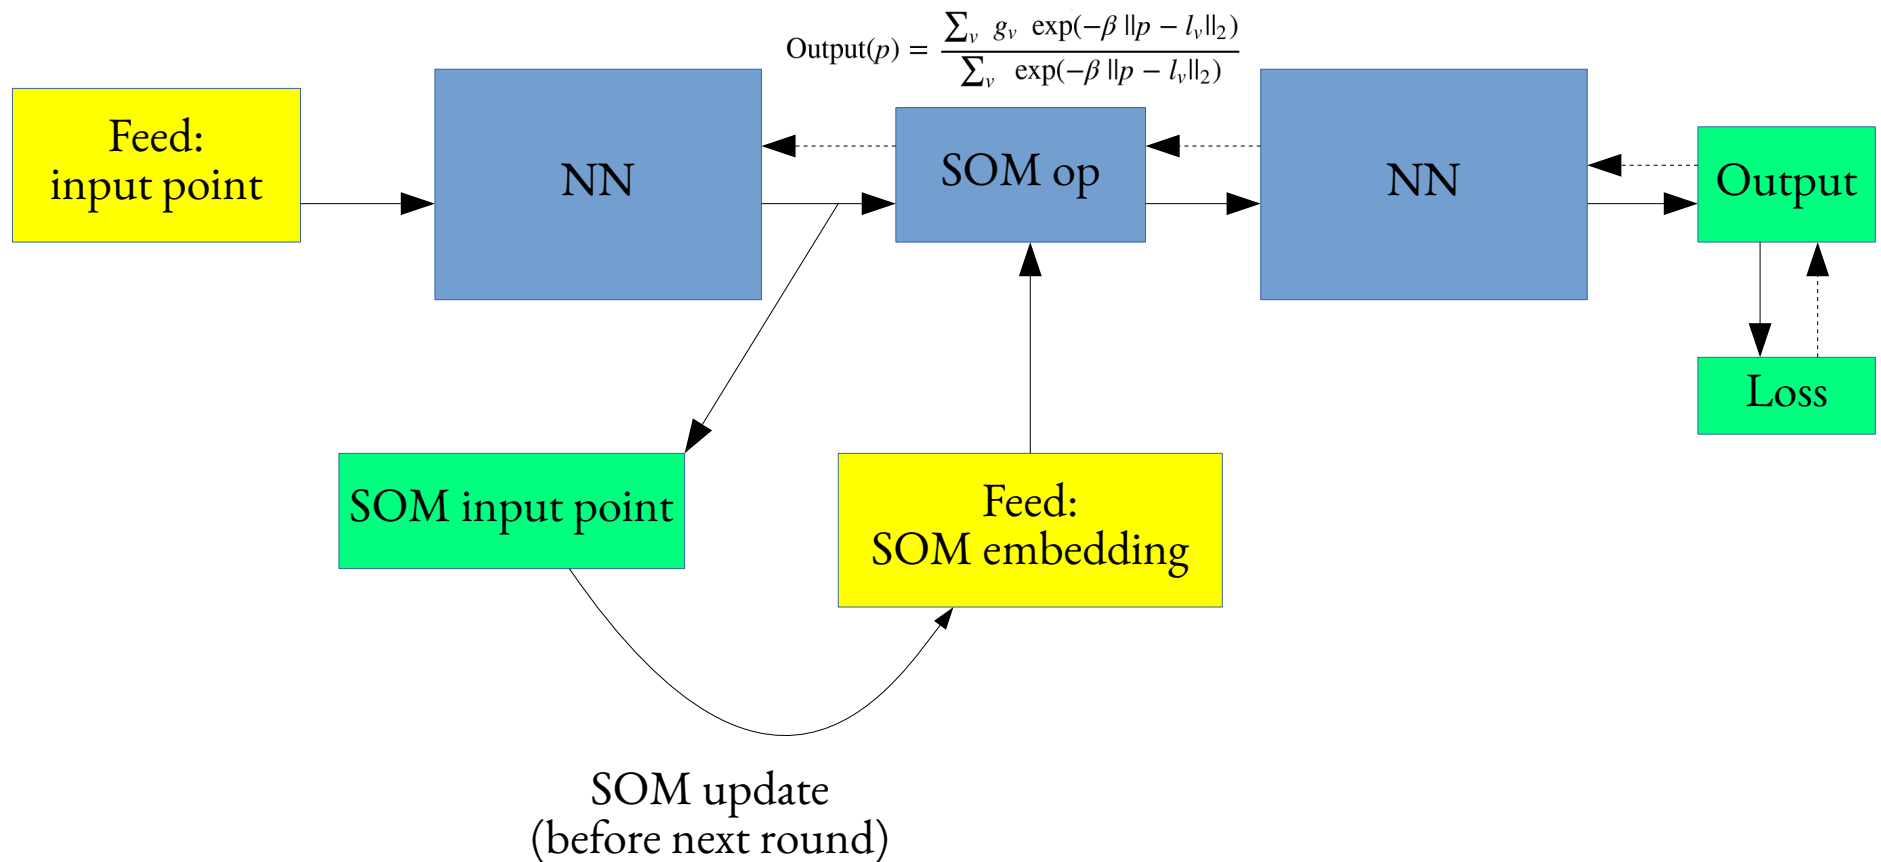
$$\text{Output}(p) = \frac{\sum_v g_v \exp(-\beta \|p - l_v\|_2)}{\sum_v \exp(-\beta \|p - l_v\|_2)}$$

Implemented as TensorFlow op,
with embedding as TF input.


TensorFlow structure




TensorFlow structure




Testing: spiral classification task



Results: pure NN





(Stochastic gradient descent)




~65% accuracy



Results: SOM preprocessing



~85% accuracy



Results: interposed SOM



~55% accuracy

Q&A